QUESTIONS

How should retrieval failures be evaluated?

Define judgments, query groups and baselines before extending the toy corpus.

Question

Which retrieval failures arise from missing vocabulary, tokenization, ranking, or incomplete relevance judgments?

Current evidence

Retrieval Lab has deterministic cosine ranking, explicit binary relevance metrics and a tiny authored corpus. Hand-computed tests validate the implementation. They do not establish quality on a broader query distribution.

Smallest useful next study

Create a licensed corpus manifest and an independently judged query set. Specify query groups and a frozen held-out set before choosing preprocessing variants. Compare the existing TF-IDF baseline with one declared alternative, retain raw rankings and report metric denominators and failure categories.

Decisions still open

Corpus/license, relevance adjudication, multilingual coverage, query count, uncertainty method and compute budget are not selected yet. No embedding-provider access or paid evaluation is assumed.

References

- scikit-learn: Text feature extraction ↗
- scikit-learn: Model evaluation ↗