QUESTIONS

Choose the next real project

A small decision record for extending the working examples.

Candidate directions

Existing implementation	Next question	Evidence needed before expansion
Linear Model Lab	Does regularization improve an ill-conditioned fit?	Fixed split, baseline, conditioning and a declared parameter sweep
Retrieval Lab	Which failures come from tokenization versus ranking?	Judged corpus, query groups, corpus/license provenance and stable metrics
Attention Lab	Do autodiff and manual gradients agree?	Independent gradient reference, tolerances, shape/mask checks and memory budget

These are possible next projects, not selected research conclusions. Keep expensive training and hosted model calls behind an explicit experiment plan with a compute budget.